

Adversarial Variational Optimization of Non-Differentiable Simulators

Gilles Louppe¹ and Kyle Cranmer¹

¹New York University

Complex computer simulators are increasingly used across fields of science as generative models tying parameters of an underlying theory to experimental observations. Inference in this setup is often difficult, as simulators rarely admit a tractable density or likelihood function. In this note, we develop Adversarial Variational Optimization (AVO), a likelihood-free inference algorithm for fitting a non-differentiable generative model to continuous or discrete data. We adapt the training procedure of generative adversarial networks by replacing the differentiable generative network with a domain-specific simulator. We solve the resulting non-differentiable minimax problem by minimizing variational upper bounds of the two adversarial objectives. Effectively, the procedure results in learning an arbitrarily tight proposal distribution over simulator parameters, such that the corresponding marginal distribution of the generated data matches the observations. We present results of the method with simulators producing both discrete and continuous data.

I. INTRODUCTION

In many fields of science such as particle physics, epidemiology, and population genetics, computer simulators are used to describe complex data generation processes. These simulators relate observations \mathbf{x} to the parameters $\boldsymbol{\theta}$ of an underlying theory or mechanistic model. In most cases, these simulators are specified as procedural implementations of forward, stochastic processes involving latent variables \mathbf{z} . Rarely do these simulators admit a tractable density (or likelihood) $p(\mathbf{x}|\boldsymbol{\theta})$. The prevalence and significance of this problem has motivated an active research effort in so-called *likelihood-free inference* algorithms such as Approximate Bayesian Computation (ABC) [1–6]. Often the simulation code involves control-flow that implies the dependence of \mathbf{x} on \mathbf{z} is non-differentiable, making inference even more difficult and precludes approaches such as Ref. [7].

In parallel, with the introduction of variational autoencoders [8] and generative adversarial networks [9], there has been a vibrant research program around implicit generative models based on neural networks [10]. While these implicit generative networks also do not admit a tractable density, they are differentiable.

Generative models based on neural networks are highly parametrized and the model parameters have no obvious interpretation. In contrast, scientific simulators can be thought of as highly regularized generative models as they typically have relatively few parameters that are endowed with some level of interpretation. In this setting, inference on the model parameters $\boldsymbol{\theta}$ is often of more interest than the latent variables \mathbf{z} .

In this note, we develop an unsupervised learning algorithm for the point estimation of the parameters $\boldsymbol{\theta}$ of a non-differentiable, implicit generative model. We adapt the adversarial training procedure of generative adversarial networks [9] by replacing the implicit generative network with a domain-based scientific simulator, and solve the resulting non-differentiable minimax problem by minimizing variational upper bounds [11, 12] of the adversarial objectives. The procedure results in learning

an arbitrarily tight proposal distribution over simulator parameters, such that the corresponding marginal distribution of the generated data matches the observations.

II. PROBLEM STATEMENT

We consider a family of parameterized densities $p(\mathbf{x}|\boldsymbol{\theta})$ defined implicitly through the simulation of a stochastic generative process, where $\mathbf{x} \in \mathbb{R}^d$ is the data and $\boldsymbol{\theta}$ are the parameters of interest. The simulation may involve some complicated latent process where $\mathbf{z} \in \mathcal{Z}$ is a latent variable providing an external source of randomness. Unlike implicit generative models defined by neural networks, we do not assume \mathbf{z} to be a fixed-size vector with a simple density. Instead, the dimensionality of \mathbf{z} and the nature of its components (uniform, normal, discrete, continuous, etc.) are inherited from the control flow of the simulation code and may depend on $\boldsymbol{\theta}$ in some intricate way. Moreover, the dimensionality of \mathbf{z} may be much larger than the dimensionality of \mathbf{x} .

We assume that the stochastic generative process that defines $p(\mathbf{x}|\boldsymbol{\theta})$ is specified through a non-differentiable deterministic function $g(\cdot; \boldsymbol{\theta}) : \mathcal{Z} \rightarrow \mathbb{R}^d$. Operationally,

$$\mathbf{x} \sim p(\mathbf{x}|\boldsymbol{\theta}) \equiv \mathbf{z} \sim p(\mathbf{z}|\boldsymbol{\theta}), \mathbf{x} = g(\mathbf{z}; \boldsymbol{\theta}) \quad (1)$$

such that the density $p(\mathbf{x}|\boldsymbol{\theta})$ can be rewritten as

$$p(\mathbf{x}|\boldsymbol{\theta}) = \int_{\{\mathbf{z}: g(\mathbf{z}; \boldsymbol{\theta}) = \mathbf{x}\}} p(\mathbf{z}|\boldsymbol{\theta}) \mu(d\mathbf{z}), \quad (2)$$

where μ is a probability measure.

Given some observed data $\{\mathbf{x}_i | i = 1, \dots, N\}$ drawn from the (unknown) true distribution $p_r(\mathbf{x})$, our goal is to estimate the parameters $\boldsymbol{\theta}^*$ that minimize the divergence between $p_r(\mathbf{x})$ and the implicit model $p(\mathbf{x}|\boldsymbol{\theta})$. That is,

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta} \in \Theta} \rho(p_r(\mathbf{x}), p(\mathbf{x}|\boldsymbol{\theta})), \quad (3)$$

where ρ is some distance or divergence. For many choices of divergence ρ , this is asymptotically equivalent to a maximum likelihood estimate.

III. BACKGROUND

A. Generative adversarial networks

Generative adversarial networks (GANs) were first proposed by [9] as a way to build an implicit generative model capable of producing samples from random noise \mathbf{z} . More specifically, a generative model $g(\cdot; \boldsymbol{\theta})$ is pit against an adversarial classifier $d(\cdot; \boldsymbol{\phi}) : \mathbb{R}^d \rightarrow [0, 1]$ with parameters $\boldsymbol{\phi}$ and whose antagonistic objective is to recognize real data \mathbf{x} from generated data $\tilde{\mathbf{x}} = g(\mathbf{z}; \boldsymbol{\theta})$. Both models g and d are trained simultaneously, in such a way that g learns to fool its adversary d (which happens when g produces samples comparable to the observed data), while d continuously adapts to changes in g . When d is trained to optimality before each parameter update of the generator, it can be shown that the original adversarial learning procedure [9] amounts to minimizing the Jensen-Shannon divergence $\text{JSD}(p_r(\mathbf{x}) \parallel p(\mathbf{x}|\boldsymbol{\theta}))$ between $p_r(\mathbf{x})$ and $p(\mathbf{x}|\boldsymbol{\theta})$.

As thoroughly explored in [13], GANs remain remarkably difficult to train because of vanishing gradients as d saturates, or because of unreliable updates when the training procedure is relaxed. As a remedy, Wasserstein GANs [14] reformulate the adversarial setup in order to minimize the Wasserstein-1 distance $W(p_r(\mathbf{x}), p(\mathbf{x}|\boldsymbol{\theta}))$ by replacing the adversarial classifier with a 1-Lipschitz adversarial critic $d(\cdot; \boldsymbol{\phi}) : \mathbb{R}^d \rightarrow \mathbb{R}$.

$$W(p_r(\mathbf{x}), p(\mathbf{x}|\boldsymbol{\theta})) = \sup_{\hat{d}} \mathbb{E}_{p(\tilde{\mathbf{x}}|\boldsymbol{\theta})}[d(\tilde{\mathbf{x}})] - \mathbb{E}_{p_r(\mathbf{x})}[d(\mathbf{x})] \quad (4)$$

Under the WGAN-GP formulation of [15] for stabilizing the optimization procedure, training d and g results in alternating gradient updates on $\boldsymbol{\phi}$ and $\boldsymbol{\theta}$ in order to respectively minimize

$$\mathcal{L}_d = W(p_r, p_\theta) + \lambda \mathbb{E}_{\tilde{\mathbf{x}} \sim p(\tilde{\mathbf{x}})} [(\|\nabla_{\tilde{\mathbf{x}}} d(\tilde{\mathbf{x}}; \boldsymbol{\phi})\|_2 - 1)^2] \quad (5)$$

$$\mathcal{L}_g = -W(p_r, p_\theta) \quad (6)$$

where $\hat{\mathbf{x}} := \epsilon \mathbf{x} + (1 - \epsilon) \tilde{\mathbf{x}}$, for $\epsilon \sim U[0, 1]$, $\mathbf{x} \sim p_r(\mathbf{x})$ and $\tilde{\mathbf{x}} \sim p(\mathbf{x}|\boldsymbol{\theta})$.

B. Variational optimization

Variational optimization [12, 16] and evolution strategies [11] are general optimization techniques that can be used to form a differentiable bound on the optima of a non-differentiable function. Given a function f to minimize, these techniques are based on the simple fact that

$$\min_{\boldsymbol{\theta} \in \Theta} f(\boldsymbol{\theta}) \leq \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})}[f(\boldsymbol{\theta})] = U(\boldsymbol{\psi}), \quad (7)$$

where $q(\boldsymbol{\theta}|\boldsymbol{\psi})$ is a proposal distribution with parameters $\boldsymbol{\psi}$ over input values $\boldsymbol{\theta}$. That is, the minimum of a set of function values is always less than or equal to any of their average. Provided that the proposal is flexible enough,

the parameters $\boldsymbol{\psi}$ can be updated to place its mass arbitrarily tight around the optimum $\boldsymbol{\theta}^* = \min_{\boldsymbol{\theta} \in \Theta} f(\boldsymbol{\theta})$.

Under mild restrictions outlined in [12], the bound $U(\boldsymbol{\psi})$ is differentiable with respect to $\boldsymbol{\psi}$, and using the log-likelihood trick it can be rewritten as:

$$\begin{aligned} \nabla_{\boldsymbol{\psi}} U(\boldsymbol{\psi}) &= \nabla_{\boldsymbol{\psi}} \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})}[f(\boldsymbol{\theta})] \\ &= \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})}[f(\boldsymbol{\theta}) \nabla_{\boldsymbol{\psi}} \log q(\boldsymbol{\theta}|\boldsymbol{\psi})] \end{aligned} \quad (8)$$

Effectively, this means that provided that the score function $\nabla_{\boldsymbol{\psi}} \log q(\boldsymbol{\theta}|\boldsymbol{\psi})$ of the proposal is known and that one can evaluate $f(\boldsymbol{\theta})$ for any $\boldsymbol{\theta}$, then one can construct empirical estimates of Eqn. 8, which can in turn be used to minimize $U(\boldsymbol{\psi})$ with stochastic gradient descent (or a variant thereof, like Adam [17] or the Natural Evolution Strategy algorithm [11], for scaling invariance and robustness to noisy gradients).

IV. ADVERSARIAL VARIATIONAL OPTIMIZATION

The alternating stochastic gradient descent on \mathcal{L}_d and \mathcal{L}_g in GANs (Section III A) inherently assumes that the generator g is a differentiable function. In the setting where we are interested in estimating the parameters of a fixed non-differentiable simulator (Section II), rather than in learning the generative model itself, gradients $\nabla_{\boldsymbol{\theta}} g$ either do not exist or are inaccessible. As a result, gradients $\nabla_{\boldsymbol{\theta}} \mathcal{L}_g$ cannot be constructed and the optimization procedure cannot be carried out.

In this work, we propose to rely on variational optimization to minimize \mathcal{L}_d and \mathcal{L}_g , thereby bypassing the non-differentiability of g . More specifically, we consider a proposal distribution $q(\boldsymbol{\theta}|\boldsymbol{\psi})$ over the parameters of g and $p(\mathbf{x}|\boldsymbol{\theta})$ and alternately minimize the variational upper bounds

$$U_d = \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})}[\mathcal{L}_d] \quad (9)$$

$$U_g = \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})}[\mathcal{L}_g] \quad (10)$$

respectively over $\boldsymbol{\phi}$ and $\boldsymbol{\psi}$. When updating $\boldsymbol{\phi}$, unbiased estimates of $\nabla_{\boldsymbol{\phi}} U_d$ can be obtained by directly evaluating the gradient of U_d over mini-batches of real and generated data, as is ordinarily done in stochastic gradient descent. When updating $\boldsymbol{\psi}$, $\nabla_{\boldsymbol{\psi}} U_g$ can be estimated as described in the previous section. That is,

$$\nabla_{\boldsymbol{\psi}} U_g = \mathbb{E}_{\substack{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi}) \\ \tilde{\mathbf{x}} \sim p(\tilde{\mathbf{x}}|\boldsymbol{\theta})}}[-d(\tilde{\mathbf{x}}; \boldsymbol{\phi}) \nabla_{\boldsymbol{\psi}} \log q(\boldsymbol{\theta}|\boldsymbol{\psi})], \quad (11)$$

which we can approximate with mini-batches of generated data

$$\nabla_{\boldsymbol{\psi}} U_g \approx \frac{1}{M} \sum_{m=1}^M -d(g(\mathbf{z}_m; \boldsymbol{\theta}_m); \boldsymbol{\phi}) \nabla_{\boldsymbol{\psi}} \log q(\boldsymbol{\theta}_m|\boldsymbol{\psi}) \quad (12)$$

for $\boldsymbol{\theta}_m \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})$ and $\mathbf{z}_m \sim p(\mathbf{z}|\boldsymbol{\theta}_m)$. For completeness, Algorithm 1 outlines the proposed adversarial variational optimization (AVO) procedure, as built on top of WGAN-GP.

Algorithm 1 Adversarial variational optimization (AVO).

Inputs: observed data $\{\mathbf{x}_i \sim p_r(\mathbf{x})\}_{i=1}^N$, simulator g .

Outputs: proposal distribution $q(\boldsymbol{\theta}|\boldsymbol{\psi})$, such that $q(\mathbf{x}|\boldsymbol{\psi}) \approx p_r(\mathbf{x})$.

Hyper-parameters: The number n_{critic} of training iterations of d ; the size M of a mini-batch; the gradient penalty coefficient λ ; the entropy penalty coefficient γ .

```

1:  $q(\boldsymbol{\theta}|\boldsymbol{\psi}) \leftarrow$  prior on  $\boldsymbol{\theta}$  (with differentiable and known density)
2: while  $\boldsymbol{\psi}$  has not converged do
3:   for  $i = 1$  to  $n_{\text{critic}}$  do ▷ Update  $d$ 
4:     Sample a mini-batch  $\{\mathbf{x}_m \sim p_r(\mathbf{x}), \boldsymbol{\theta}_m \sim q(\boldsymbol{\theta}|\boldsymbol{\psi}), \mathbf{z}_m \sim p(\mathbf{z}|\boldsymbol{\theta}_m), \epsilon_m \sim U[0, 1]\}_{m=1}^M$ .
5:     for  $m = 1$  to  $M$  do
6:        $\tilde{\mathbf{x}}_m \leftarrow g(\mathbf{z}_m; \boldsymbol{\theta}_m)$ 
7:        $\hat{\mathbf{x}}_m \leftarrow \epsilon_m \mathbf{x}_m + (1 - \epsilon_m) \tilde{\mathbf{x}}_m$ 
8:        $U_d^{(m)} \leftarrow d(\tilde{\mathbf{x}}_m; \boldsymbol{\phi}) - d(\mathbf{x}_m; \boldsymbol{\phi}) + \lambda(\|\nabla_{\tilde{\mathbf{x}}_m} d(\tilde{\mathbf{x}}_m; \boldsymbol{\phi})\|_2 - 1)^2$ 
9:     end for
10:     $\boldsymbol{\phi} \leftarrow \text{Adam}(\nabla_{\boldsymbol{\phi}} \frac{1}{M} \sum_{m=1}^M U_d^{(m)})$ 
11:  end for
12:  Sample a mini-batch  $\{\boldsymbol{\theta}_m \sim q(\boldsymbol{\theta}|\boldsymbol{\psi}), \mathbf{z}_m \sim p(\mathbf{z}|\boldsymbol{\theta}_m)\}_{m=1}^M$ . ▷ Update  $q(\boldsymbol{\theta}|\boldsymbol{\psi})$ 
13:   $\nabla_{\boldsymbol{\psi}} U_g \leftarrow \frac{1}{M} \sum_{m=1}^M -d(g(\mathbf{z}_m; \boldsymbol{\theta}_m)) \nabla_{\boldsymbol{\psi}} \log q_{\boldsymbol{\psi}}(\boldsymbol{\theta}_m)$ 
14:   $\nabla_{\boldsymbol{\psi}} H(q_{\boldsymbol{\psi}}) \leftarrow \frac{1}{M} \sum_{m=1}^M \nabla_{\boldsymbol{\psi}} q_{\boldsymbol{\psi}}(\boldsymbol{\theta}_m) \log q_{\boldsymbol{\psi}}(\boldsymbol{\theta}_m)$ 
15:   $\boldsymbol{\psi} \leftarrow \text{Adam}(\nabla_{\boldsymbol{\psi}} U_g + \gamma \nabla_{\boldsymbol{\psi}} H(q_{\boldsymbol{\psi}}))$ 
16: end while

```

Algorithm 1 represents the simplest version of AVO; however, the variance of the noisy estimator of the gradients may be too large to be useful in many problems. Variance reduction techniques such as Rao-Blackwellization and Control Variates can be applied here as they are in black box variational inference [19].

A. Adversarial variational inference through model criticism

The variational objectives 9-10 have the effect of replacing the modeled data distribution of Eqn. 1 with the marginal distribution parameterized by $\boldsymbol{\psi}$

$$q(\mathbf{x}|\boldsymbol{\psi}) = \int p(\mathbf{x}|\boldsymbol{\theta}) q(\boldsymbol{\theta}|\boldsymbol{\psi}) d\boldsymbol{\theta}. \quad (13)$$

We can think of $p(\mathbf{x}|\boldsymbol{\psi})$ as a *variational program* as described in [29], though more complicated than a simple reparameterization of normally distributed noise \mathbf{z} through a differentiable function $g(\mathbf{z}, \boldsymbol{\psi})$. In our case, the variational program is a marginalized, non-differentiable simulator. Its density is intractable; nevertheless, it can generate samples for \mathbf{x} and is differentiable with respect to $\boldsymbol{\psi}$ and $\boldsymbol{\phi}$. Operationally, we sample from this marginal model via

$$\mathbf{x} \sim q(\mathbf{x}|\boldsymbol{\psi}) \equiv \boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi}), \mathbf{z} \sim p(\mathbf{z}|\boldsymbol{\theta}), \mathbf{x} = g(\mathbf{z}; \boldsymbol{\theta}). \quad (14)$$

Effectively, we are optimizing variational posterior $q(\boldsymbol{\theta}|\boldsymbol{\psi})$ by minimizing the divergence of the corresponding posterior predictive distribution $q(\mathbf{x}|\boldsymbol{\psi})$ and the data distribution $p_r(\mathbf{x})$. Consequently, the proposed algorithm is a form of variational inference through model criticism [18], which we denote AVI-MC.

B. Adversarial variational learning

If instead, one is interested in point estimates $\boldsymbol{\theta}^*$, then we augment Eqn. 10 with a regularization term corresponding to the differential entropy H of the proposal distribution $q(\boldsymbol{\theta}|\boldsymbol{\psi})$. That is,

$$U_g = \mathbb{E}_{\boldsymbol{\theta} \sim q(\boldsymbol{\theta}|\boldsymbol{\psi})} [\mathcal{L}_g] + \gamma H(q(\boldsymbol{\theta}|\boldsymbol{\psi})) \quad (15)$$

where $\gamma \in \mathbb{R}^+$ is a hyper-parameter controlling the trade-off between the generator objective and the tightness of the proposal distribution. For small values of γ , proposal distributions with large entropy are not penalized, which results in learning a smeared variation of the original simulator corresponding to the posterior predictive model. On the other hand, for large values of γ , the procedure is encouraged to fit a proposal distribution with low entropy, which has the effect of concentrating its density tightly around one or a few $\boldsymbol{\theta}$ values. Very large penalties may eventually make the optimization unstable, as the variance of $\nabla_{\boldsymbol{\psi}} \log q(\boldsymbol{\theta}_m|\boldsymbol{\psi})$ typically increases as the entropy of the proposal decreases. [KC: maybe, but not obvious why]

[KC: Discuss hierarchical models and iid data \mathbf{x} . Factor of N/M for AVI-MC?]




FIG. 1. Discrete Poisson model with unknown mean. (*Top left*) Proposal distributions $q(\theta|\psi)$ after adversarial variational optimization. For both $\gamma = 0$ and $\gamma = 5$, the distributions correctly concentrate their density around the true value $\log(\lambda^*)$. Penalizing the entropy of the proposal distribution ($\gamma = 5$) results in a tighter density. (*Top right*) Model distributions $q(\mathbf{x}|\psi)$ after training. This plot shows that the resulting parameterizations of the simulator closely reproduce the true distribution. (*Bottom*) Empirical estimates of the variational upper bound U_d as optimization progresses.

V. EXPERIMENTS

A. Univariate discrete data

As a first illustrative experiment, we evaluate inference for a discrete Poisson distribution with unknown mean λ . We artificially consider the distribution as a parameterized simulator, from which we can only generate data.

The observed data is sampled from a Poisson with mean $\lambda^* = 7$. Algorithm 1 is run for 300 epochs with mini-batches of size $M = 64$ and the following configuration. For the critic d , we use a 3-layer MLP with 10 hidden nodes per layer and ReLU activations. At each epoch, Adam is run for $n_{\text{critic}} = 100$ iterations with a step size $\alpha = 0.01$, decay rates $\beta_1 = \beta_2 = 0.5$ and its inner first and second moment vectors reset at each outer epoch in order to avoid building momentum in staled directions. For estimating λ^* , we parameterize θ as $\log(\lambda)$ and use a univariate Gaussian proposal distribution $q(\theta|\psi)$ initialized with a mean at $\log(5)$ and unit variance. At each epoch, parameters ψ are updated by taking one Adam step, with $\alpha = 0.01$ and $\beta_1 = \beta_2 = 0.5$. The gradient penalty coefficient is set to $\lambda = 0.025$, and the entropy penalty is evaluated at both $\gamma = 0$ and $\gamma = 5$.

The top left plot in Figure 1 illustrates the resulting proposal distributions $q(\theta|\psi)$ after AVO. For both $\gamma = 0$ and $\gamma = 5$, the proposal distributions correctly concentrate their density around the true parameter value $\log(\lambda^*) = 1.94$. Under the effect of the positive entropy penalty $H(q(\theta|\psi))$, the proposal distribution for $\gamma = 5$ concentrates its mass more tightly, yielding in this case a more precise inference. The top right plot compares the model distributions to the true distribution. As theoretically expected from adversarial training, we see that the resulting distributions closely match the true distribution, with in this case visually slightly better results for the penalized model. The bottom plot of Figure 1 shows empirical estimates of $-U_d$ with respect to the epoch number. For both $\gamma = 0$ and $\gamma = 5$, the curves quickly fall towards 0, which indicates that $\mathbb{E}_{\tilde{\mathbf{x}} \sim p(\mathbf{x}|\theta)}[d(\tilde{\mathbf{x}}; \phi)] \approx \mathbb{E}_{\mathbf{x} \sim p_r(\mathbf{x})}[d(\mathbf{x}; \phi)]$ and that the critic cannot distinguish between true and model data. Despite the discreteness and the non-differentiability of the underlying generator, this confirms that inference with adversarial variational optimization works.

B. Multidimensional continuous data

As a second toy example, we evaluate parameter inference for a generator producing 5-dimensional continuous data, as originally specified in Section 4.2. of [6]. More specifically, we consider the following generative process:

- $\mathbf{z} = (z_0, z_1, z_2, z_3, z_4)$, such that $z_0 \sim \mathcal{N}(\mu = \alpha, \sigma = 1)$, $z_1 \sim \mathcal{N}(\mu = \beta, \sigma = 3)$, $z_2 \sim \text{Mixture}(\frac{1}{2}\mathcal{N}(\mu = -2, \sigma = 1), \frac{1}{2}\mathcal{N}(\mu = 2, \sigma = 0.5))$, $z_3 \sim \text{Exponential}(\lambda = 3)$, and $z_4 \sim \text{Exponential}(\lambda = 0.5)$;
- $\mathbf{x} = R\mathbf{z}$, where R is a fixed semi-positive definite 5×5 matrix defining a fixed projection of \mathbf{z} into the observed space.

We consider observed data generated at the nominal values $\theta^* = (\alpha^* = 1, \beta^* = -1)$. The simulator parameters are modeled with a factored Gaussian proposal distribution $q(\theta|\psi) = q(\alpha|\psi)q(\beta|\psi)$, where each component is initialized with zero mean and unit variance. Hyperparameters are set to $M = 64$, $n_{\text{critic}} = 100$, $\lambda = 0.025$, $\gamma = 10$ and Adam configured with $\alpha = 0.01$, $\beta_1 = 0.9$ and $\beta_2 = 0.999$. The network architecture of the critic is the same as in the previous example.

Starting with a proposal distribution $q(\theta|\psi)$ largely spread over the parameter space, as illustrated in the left plot of Figure 2, inference quickly converges towards a proposal distribution whose density concentrates around the nominal values θ^* , as shown in the right plot of Figure 2. Overall, this example further illustrates and confirms the ability of adversarial variational optimization for inference with multiple parameters and multidimensional data, where reliable approximations of $p(\mathbf{x}|\theta)$ in a traditional MLE setting would otherwise be difficult to construct.




FIG. 2. Multidimensional continuous data. (Left) Density $q(\theta|\psi)$ at the beginning of the procedure, for a proposal distribution initialized with zero mean and unit variance. (Right) Density $q(\theta|\psi)$ after adversarial variational optimization ($\gamma = 10$). The proposal density correctly converges towards a distribution whose density concentrates around $\theta^* = (1, -1)$.

C. Electron-positron annihilation

As a more challenging example, we now consider a (simplified) simulator from particle physics for electron-positron collisions resulting in muon-antimuon pairs ($e^+e^- \rightarrow \mu^+\mu^-$). The simulator approximates the distribution of observed measurements $\mathbf{x} = \cos(A) \in [-1, 1]$, where A is the polar angle of the outgoing muon with respect to the originally incoming electron. This random variable is approximately distributed as

$$p(\mathbf{x}|E_{\text{beam}}, G_f) = \frac{1}{Z} [(1 + \mathbf{x}^2) + c(E_{\text{beam}}, G_f)\mathbf{x}] \quad (16)$$

where Z is a known normalization constant and c is an asymmetry coefficient function. Due to the linear term in the expression, the density $p(\mathbf{x}|E_{\text{beam}}, G_f)$ exhibits a so-called *forward-backward* asymmetry. Its size depends on the values of the parameters E_{beam} (the beam energy) and G_f (the Fermi constant) through the coefficient function c . Given the known analytic form of the probability density function, a simulator for \mathbf{x} can be implemented with rejection sampling [20]. As an arbitrary number of random drawings may be required before passing the acceptance threshold, this illustrates how the latent variable \mathbf{z} does not necessarily need to be fixed-sized vector, and how its density $p(\mathbf{z}|\theta)$ may depend on θ in some complicated way (Section II).

In this example, we consider observed data generated with $\theta^* = (E_{\text{beam}}^* = 42, G_f^* = 0.9)$. The simulator parameters are modeled with a factored Gaussian proposal distribution $q(\theta|\psi) = q(E_{\text{beam}}|\psi)q(G_f|\psi)$, where each component is respectively initialized with mean 45 and 1 and variance 1 and 0.01. Hyper-parameters are set to $M = 64$, $n_{\text{critic}} = 100$, $\lambda = 0.0025$ and Adam configured with $\alpha = 0.01$, $\beta_1 = 0.9$ and $\beta_2 = 0.999$. As for the first example, we compare entropy penalties $\gamma = 0$ and $\gamma = 5$.

The top left plot in Figure 3 illustrates the resulting proposal distributions $q(\theta|\psi)$ for $\gamma = 0$ and $\gamma = 5$ after AVO. We see that the distributions both arrive in the




FIG. 3. Electron-positron annihilation. (Top left) Proposal distributions $q(\theta|\psi)$ after adversarial variational optimization. The density of the penalized distribution ($\gamma = 5$) is here highly concentrated, resulting in the green mass near θ^* . (Top right) Model distributions $q(\mathbf{x}|\psi)$ after training. Despite the differences between their respective proposal distributions, both models closely match the observed data. (Bottom) Empirical estimates of the variational upper bound $-U_d$ as optimization progresses.

neighborhood of θ^* , with a density more highly concentrated for $\gamma = 5$ than for $\gamma = 0$. Despite these differences and the relative distance with θ^* , both models closely match the observed data, as shown in the top right plot of Figure 3, with again slightly better results for the entropy penalized model. This suggests either a relatively flat landscape around θ^* or that the observed data can in this case also be reproduced with a smeared variation $q(\mathbf{x}|\psi)$ of the simulator. Finally, the bottom plot of Figure 3 shows that for both $\gamma = 0$ and $\gamma = 5$ the variational upper bound $-U_d$ quickly fall towards 0, which indicates convergence towards a distribution that the critic cannot distinguish from the true distribution. In summary, despite the intricate relation between the observations \mathbf{x} , the latent variable \mathbf{z} and the parameters θ , this example demonstrates the efficacy of AVO for parameter inference on a black box simulator.

VI. RELATED WORKS

As reviewed in [10], likelihood-free inference is intimately tied to a class of algorithms that can be framed as density estimation-by-comparison. In most cases, these inference algorithms are formulated as an iterative two-step process where the model distribution is first com-

pared to the true data distribution and then updated to make it more comparable to the latter.

Closest to our work are procedures that rely on a classifier to estimate the discrepancy between the true and the model distributions. For example, [21] uses non linear logistic regression for fitting unnormalized differentiable statistical models, while [9] exploits an adversarial neural network for learning a differentiable implicit generative model. In the likelihood-free setup, [6, 22, 23] estimate likelihood ratios through supervised classification, which can in turn be used for parameter inference in combination with a gradient-free optimization algorithm. Similarly, [24] makes use of classification accuracy as a summary statistics for approximate Bayesian computation.

[25]

In this context, the proposed method can be considered as a direct adaptation of generative adversarial networks [9] to non-differentiable simulators. It also constitutes an approximate gradient descent alternative to likelihood-free inference based on bayesian optimization and approximated density ratios [6, 22] or to classifier ABC [24].

This work has many connections to variational inference, in which the goal is to optimize the recognition model $q(\theta|\psi)$ so that it is close to the true posterior $p(\theta|x)$. This is slightly different from variational optimization (or variational learning), where the goal is to find θ^* such that $p(x|\theta^*) \approx p_r(x)$. Typically in variational inference the model $p(x|\theta)$ admits a density, so that the likelihood is known. Thus, assuming the prior also admits a density, it is tractable to evaluate the posterior is (up to a constant factor). In Adversarial Variational Bayes [26] and, equivalently, the PC-Adv algorithm of [27]), the authors use an adversarial setup, but assume a known density $p(x|\theta)$ that is differentiable with respect to θ .

In [28], the authors consider Variational Bayes with an intractable likelihood. In that approach "The only requirement is that the intractable likelihood can be estimated unbiasedly." In the case of most simulators it is not clear how one would obtain the unbiased estimator $\hat{p}(x|\theta)$.

In [27] the author defines implicit distributions to be "probability models whose probability density function may be intractable, but there is a way to 1) sample from them exactly and/or calculate and approximate expectations under them, and 2) calculate or estimate gradients of such expectations with respect to model parameters." This notion of implicit model does not include most scientific simulators as 2) is not satisfied. The JC-Adv and JC-Den algorithm defined there applies to implicit models, but, as the author points out, is aimed at variational inference and "does not provide a direct way to learn model parameters θ ".

A. Operator Variational Optimization

Original OPVI [29]

$$\lambda^* = \inf_{\lambda} \sup_{\phi} \mathbb{E}_{z \sim q(z|\lambda)} [(O^{p,q} f_{\lambda})] \quad (17)$$

with

$$(O^{p,q} f) = \log q(z) - \log p(z|x) \forall f \in \mathcal{F} \quad (18)$$

To get into OPVI formalism

$$\psi^* = \inf_{\psi} \sup_{\phi} \mathbb{E}_{x \sim p_r(x)} [f(x)] - \mathbb{E}_{x \sim p(x|\psi)} [f(x)] \quad (19)$$

$$= \inf_{\psi} \sup_{\phi} \mathbb{E}_{x \sim p(x|\psi)} [(O^{p_r, p\psi} f_{\phi})] \quad (20)$$

with

$$(O^{p,q} f) = \left(\frac{p_r(x)}{p(x|\psi)} - 1 \right) f(x) \quad (21)$$

We can think of $p(x|\psi)$ as a *variational program* as described in [29], though more complicated than the a simple reparametrization of normally distributed noise ϵ through a differentiable function $R(\epsilon, \psi)$. In our case, the variational program is a marginalized non-differentiable simulator simulator. Nevertheless, it can generate samples for x , is differentiable with respect to ϕ , and its density is be intractable.

VII. SUMMARY

In this note, we develop a likelihood-free inference algorithm for fitting a forward non-differentiable simulator to continuous or discrete observed data. The algorithm combines adversarial training with variational optimization to minimize variational upper bounds on the otherwise non-differentiable adversarial objectives. Effectively, the procedure results in learning an arbitrarily tight proposal distribution over the simulator parameters, such that corresponding the marginal distribution of the generated data matches the observations. Preliminary results with non-differentiable discrete or continuous simulators validate the proposed method.

While the obtained results are encouraging, the complete validation of the method remains to be carried out in real conditions on a full fledged scientific simulator – which we plan to achieve for a next version of this work. In terms of method, several components need to be further investigated. First, we need to better study the interplay between the entropy penalty and the adversarial objectives. Second, we should better understand the dynamics of the optimization procedure, in particular when combined with momentum-based optimizers like Adam. Third, we need to consider whether less noisy estimates

of the gradients $\nabla_{\psi} U_g$ can be computed, in particular by exploiting the fact in the composition $d(g(\mathbf{z}; \boldsymbol{\theta}); \psi)$ the derivatives of d are known exactly.

[KC: could also use this to tune non-differentiable fast simulator to slower full simulator. Then fixed data size not an issue]

ACKNOWLEDGMENTS

We would like to thank Lukas Heinrich for helpful comments regarding the electron–positron annihilation sim-

ulation.

GL and KL are both supported through NSF ACI-1450310, additionally KC is supported through PHY-1505463 and PHY-1205376.

-
- [1] Mark A Beaumont, Wenyang Zhang, and David J Balding. Approximate bayesian computation in population genetics. *Genetics*, 162(4):2025–2035, 2002.
 - [2] Paul Marjoram, John Molitor, Vincent Plagnol, and Simon Tavaré. Markov chain monte carlo without likelihoods. *Proceedings of the National Academy of Sciences*, 100(26):15324–15328, 2003.
 - [3] Scott A Sisson, Yanan Fan, and Mark M Tanaka. Sequential monte carlo without likelihoods. *Proceedings of the National Academy of Sciences*, 104(6):1760–1765, 2007.
 - [4] Scott A Sisson and Yanan Fan. *Likelihood-free MCMC*. Chapman & Hall/CRC, New York.[839], 2011.
 - [5] Jean-Michel Marin, Pierre Pudlo, Christian P Robert, and Robin J Ryder. Approximate bayesian computational methods. *Statistics and Computing*, pages 1–14, 2012.
 - [6] Kyle Cranmer, Juan Pavez, and Gilles Louppe. Approximating likelihood ratios with calibrated discriminative classifiers. *arXiv preprint arXiv:1506.02169*, 2015, 1506.02169.
 - [7] M. M. Graham and A. J. Storkey. Asymptotically exact inference in differentiable generative models. *ArXiv e-prints*, May 2016, 1605.07826.
 - [8] Diederik P. Kingma and Max Welling. Auto-encoding variational bayes. *CoRR*, abs/1312.6114, 2013.
 - [9] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems*, pages 2672–2680, 2014.
 - [10] S. Mohamed and B. Lakshminarayanan. Learning in Implicit Generative Models. *ArXiv e-prints*, October 2016, 1610.03483.
 - [11] D. Wierstra, T. Schaul, T. Glasmachers, Y. Sun, and J. Schmidhuber. Natural Evolution Strategies. *ArXiv e-prints*, June 2011, 1106.4487.
 - [12] J. Staines and D. Barber. Variational Optimization. *ArXiv e-prints*, December 2012, 1212.4507.
 - [13] M. Arjovsky and L. Bottou. Towards Principled Methods for Training Generative Adversarial Networks. *ArXiv e-prints*, January 2017, 1701.04862.
 - [14] M. Arjovsky, S. Chintala, and L. Bottou. Wasserstein GAN. *ArXiv e-prints*, January 2017, 1701.07875.
 - [15] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville. Improved Training of Wasserstein GANs. *ArXiv e-prints*, March 2017, 1704.00028.
 - [16] J Staines and D Barber. Optimization by variational bounding. In *ESANN 2013 proceedings, 21st European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning*, pages 473–478, 2013.
 - [17] D. P. Kingma and J. Ba. Adam: A Method for Stochastic Optimization. *ArXiv e-prints*, December 2014, 1412.6980.
 - [18] George EP Box. Sampling and bayes’ inference in scientific modelling and robustness. *Journal of the Royal Statistical Society. Series A (General)*, pages 383–430, 1980.
 - [19] Rajesh Ranganath, Sean Gerrish, and David Blei. Black box variational inference. In *Artificial Intelligence and Statistics*, pages 814–822, 2014.
 - [20] John Von Neumann. 13. various techniques used in connection with random digits. 1953.
 - [21] Michael U Gutmann and Aapo Hyvärinen. Noise-contrastive estimation of unnormalized statistical models, with applications to natural image statistics. *Journal of Machine Learning Research*, 13(Feb):307–361, 2012.
 - [22] K Cranmer, J Pavez, G Louppe, and WK Brooks. Experiments using machine learning to approximate likelihood ratios for mixture models. In *Journal of Physics: Conference Series*, volume 762, page 012034. IOP Publishing, 2016.
 - [23] R. Dutta, J. Corander, S. Kaski, and M. U. Gutmann. Likelihood-free inference by ratio estimation. *ArXiv e-prints*, November 2016, 1611.10242.
 - [24] Michael U Gutmann, Ritabrata Dutta, Samuel Kaski, and Jukka Corander. Likelihood-free inference via classification. *Statistics and Computing*, pages 1–15, 2017.
 - [25] Edward Meeds, Robert Leenders, and Max Welling. Hamiltonian abc. *arXiv preprint arXiv:1503.01916*, 2015.
 - [26] Lars M. Mescheder, Sebastian Nowozin, and Andreas Geiger. Adversarial variational bayes: Unifying variational autoencoders and generative adversarial networks. *CoRR*, abs/1701.04722, 2017.
 - [27] F. Huszár. Variational Inference using Implicit Distributions. *ArXiv e-prints*, February 2017, 1702.08235.
 - [28] Minh-Ngoc Tran, David J Nott, and Robert Kohn. Variational bayes with intractable likelihood. *Journal of Computational and Graphical Statistics*, (just-accepted), 2017, 1503.08621.
 - [29] R. Ranganath, J. Alotaar, D. Tran, and D. M. Blei. Operator Variational Inference. *ArXiv e-prints*, October

2016, 1610.09033.